

FIGHTING NON-DETERMINISM IN C++ COMPILERS


Mandeep Singh Grang • Qualcomm Innovation Center, Inc., San Diego, CA



• A compiler is non-deterministic if its output differs from run to run on the same input •

1. THE PROBLEM

C++ compilers may exhibit non-deterministic behavior.



2. WHERE IS IT OBSERVED?

- i. Back-to-back runs of the same compiler.
- ii. The same compiler hosted on different operating systems.
- iii. Asserts vs non-asserts versions of the same compiler.

4. WHAT CAUSES THIS NON-DETERMINISM?



Iteration of unordered containers


```
std::unordered_set<int> S ( {-1, 0, 1} );
```

```
for (auto I : S)
    cout << I;
```

Iteration order of the unordered_set depends on hash values of elements

Output on Ubuntu Linux x86: 1 -1 0
Output on Windows x86: -1 0 1
(compiled with LLVM 6.0)

3. WHY IS IT A PROBLEM?



Hard-to-reproduce bugs



Difficult debugging



Runtime crashes



Unpredictable performance

"If you can't reproduce a bug,
you can't fix it."

- Anonymous Programmer



Hashing of pointer keys

```
int x = -1, y = 0, z = 1;
std::map<int *, int> M;
```

```
M[&x] = x;
M[&y] = y;
M[&z] = z;
```

Keys of the map are addresses which may change from run to run

```
for (auto &I : M)
    cout << I.second;
```

Output on Ubuntu Linux x86: 1 -1 0
Output on Windows x86: 1 0 -1
(compiled with LLVM 6.0)



Use of non-stable sort functions

```
using IntPair = std::pair<int, int>;
std::vector<IntPair> V = {{0, 1}, {0, 2}};
```

```
std::sort(V.begin(), V.end(),
[] (IntPair x, IntPair y) {
    return x.first < y.first; });
```

```
for (auto &I : V)
    cout << I.second;
```

Relative ordering of keys with same values is non-deterministic

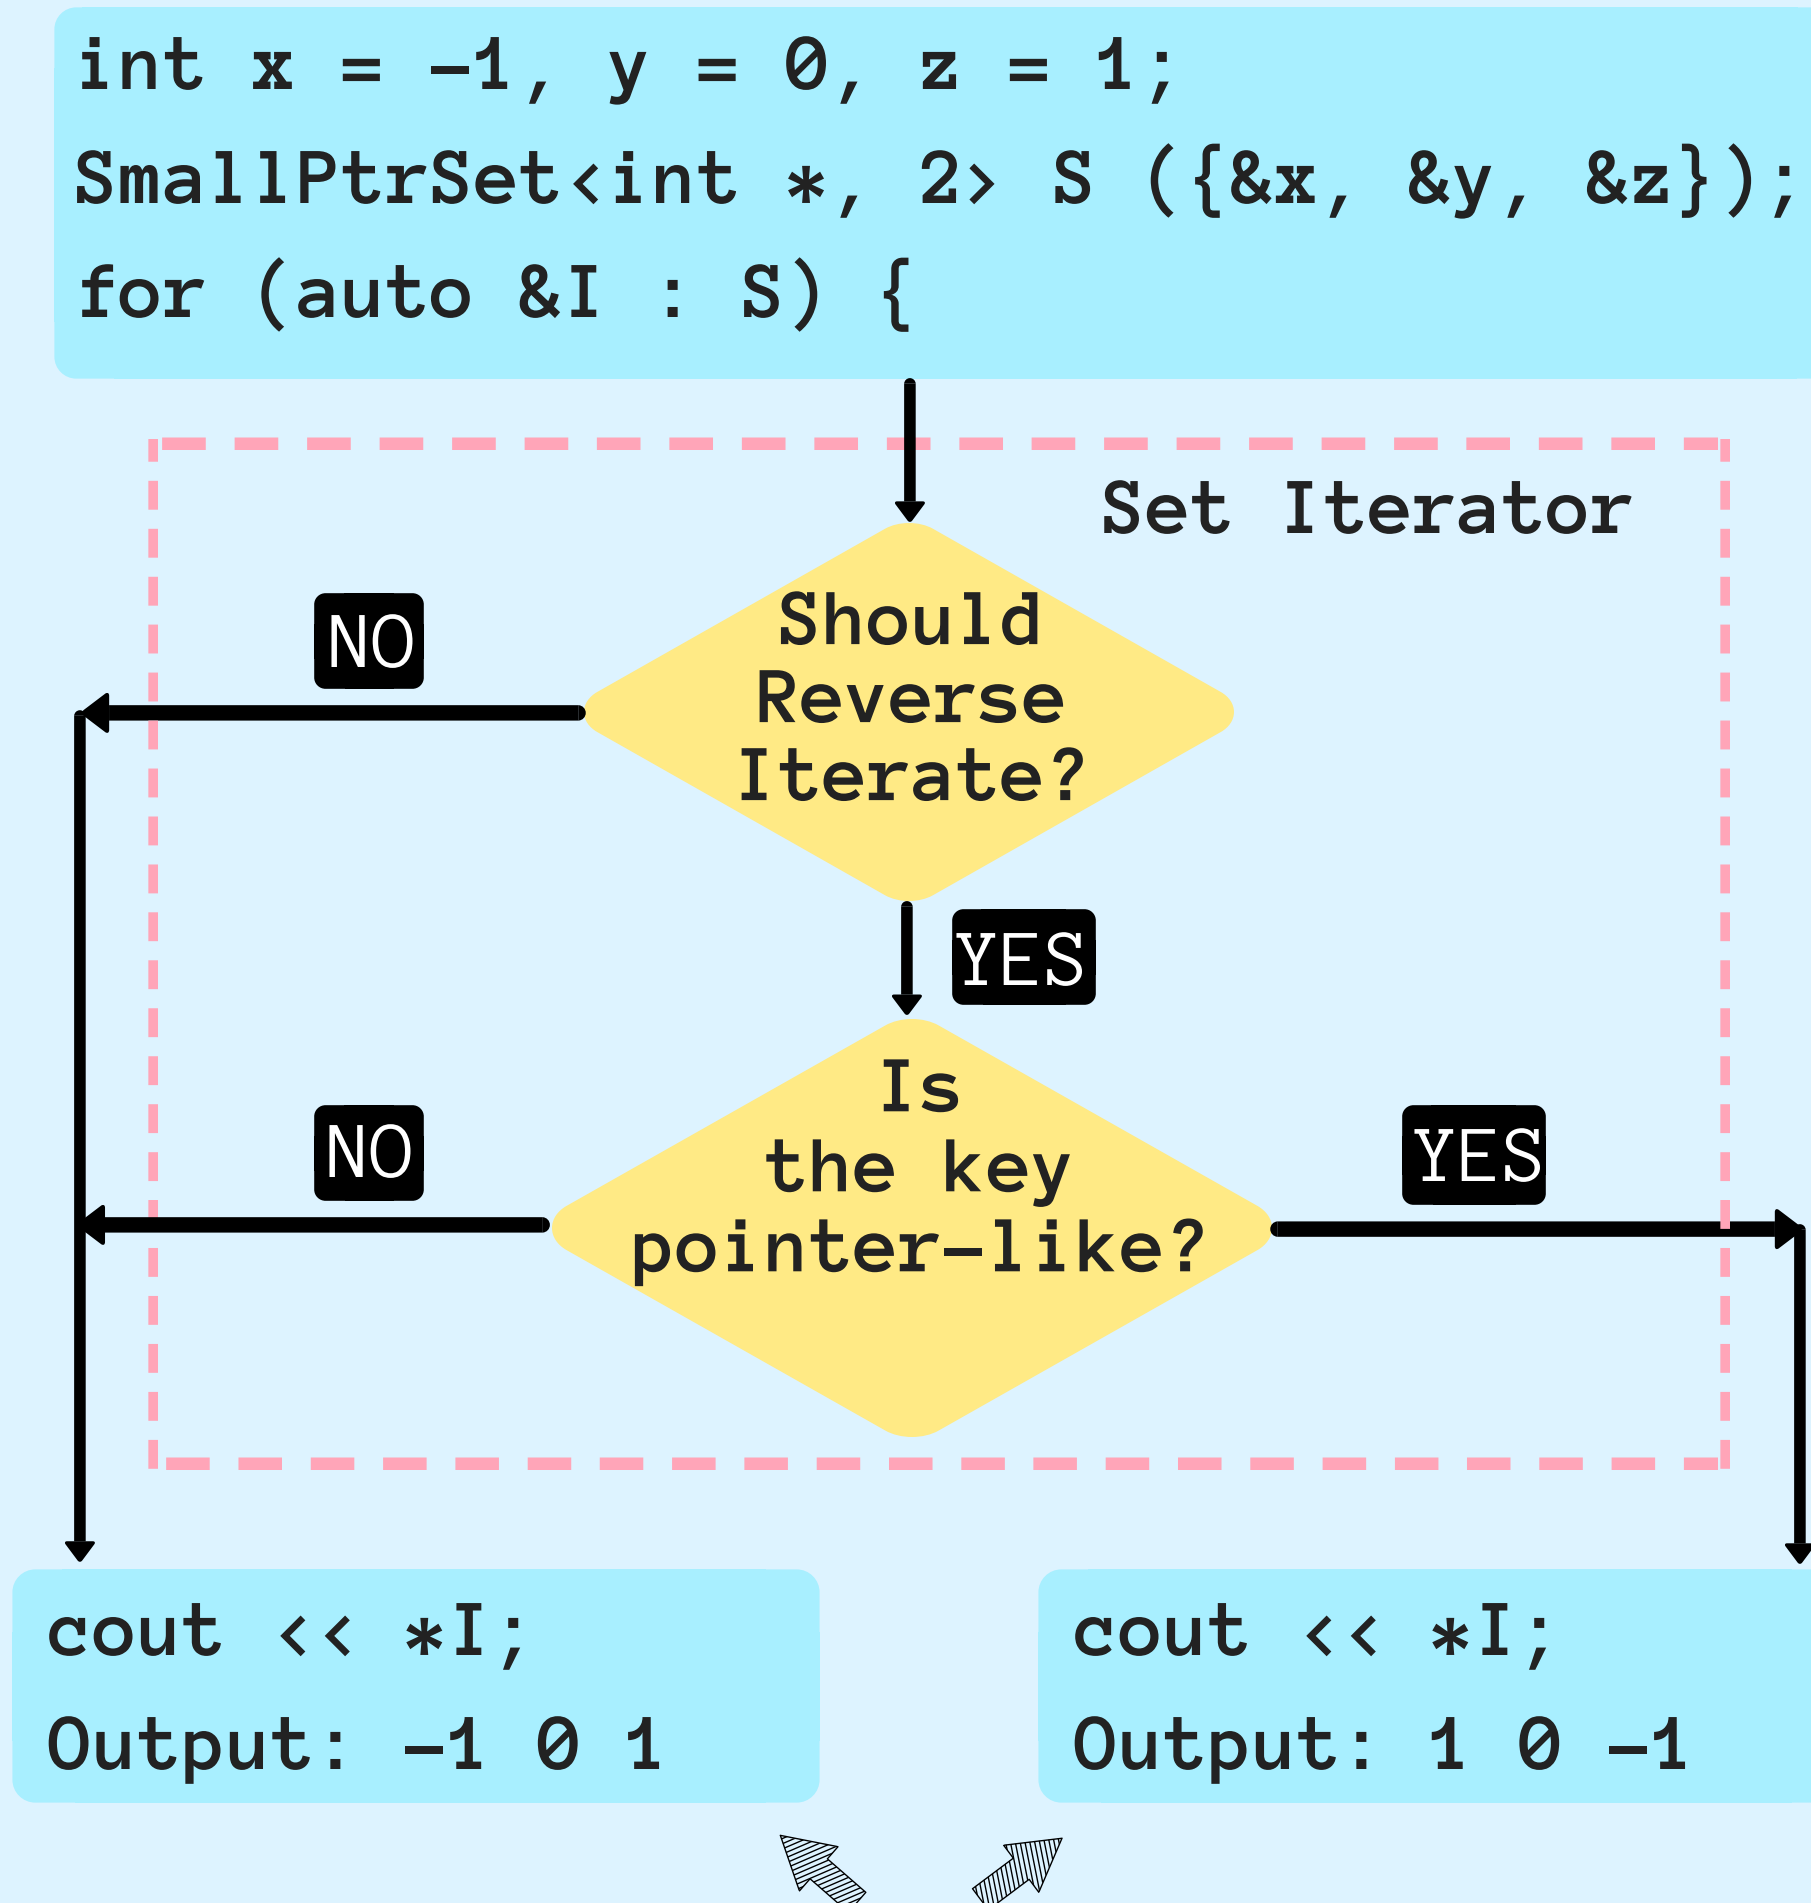
Possible outputs: 1 2 or 2 1



5. HOW DOES LLVM DETECT NON-DETERMINISTIC ITERATION ORDER?

- Added a reverse iteration mode for unordered containers.
- Compare the output of reverse iteration mode with forward iteration mode.
- Bugs uncovered and fixed:

42

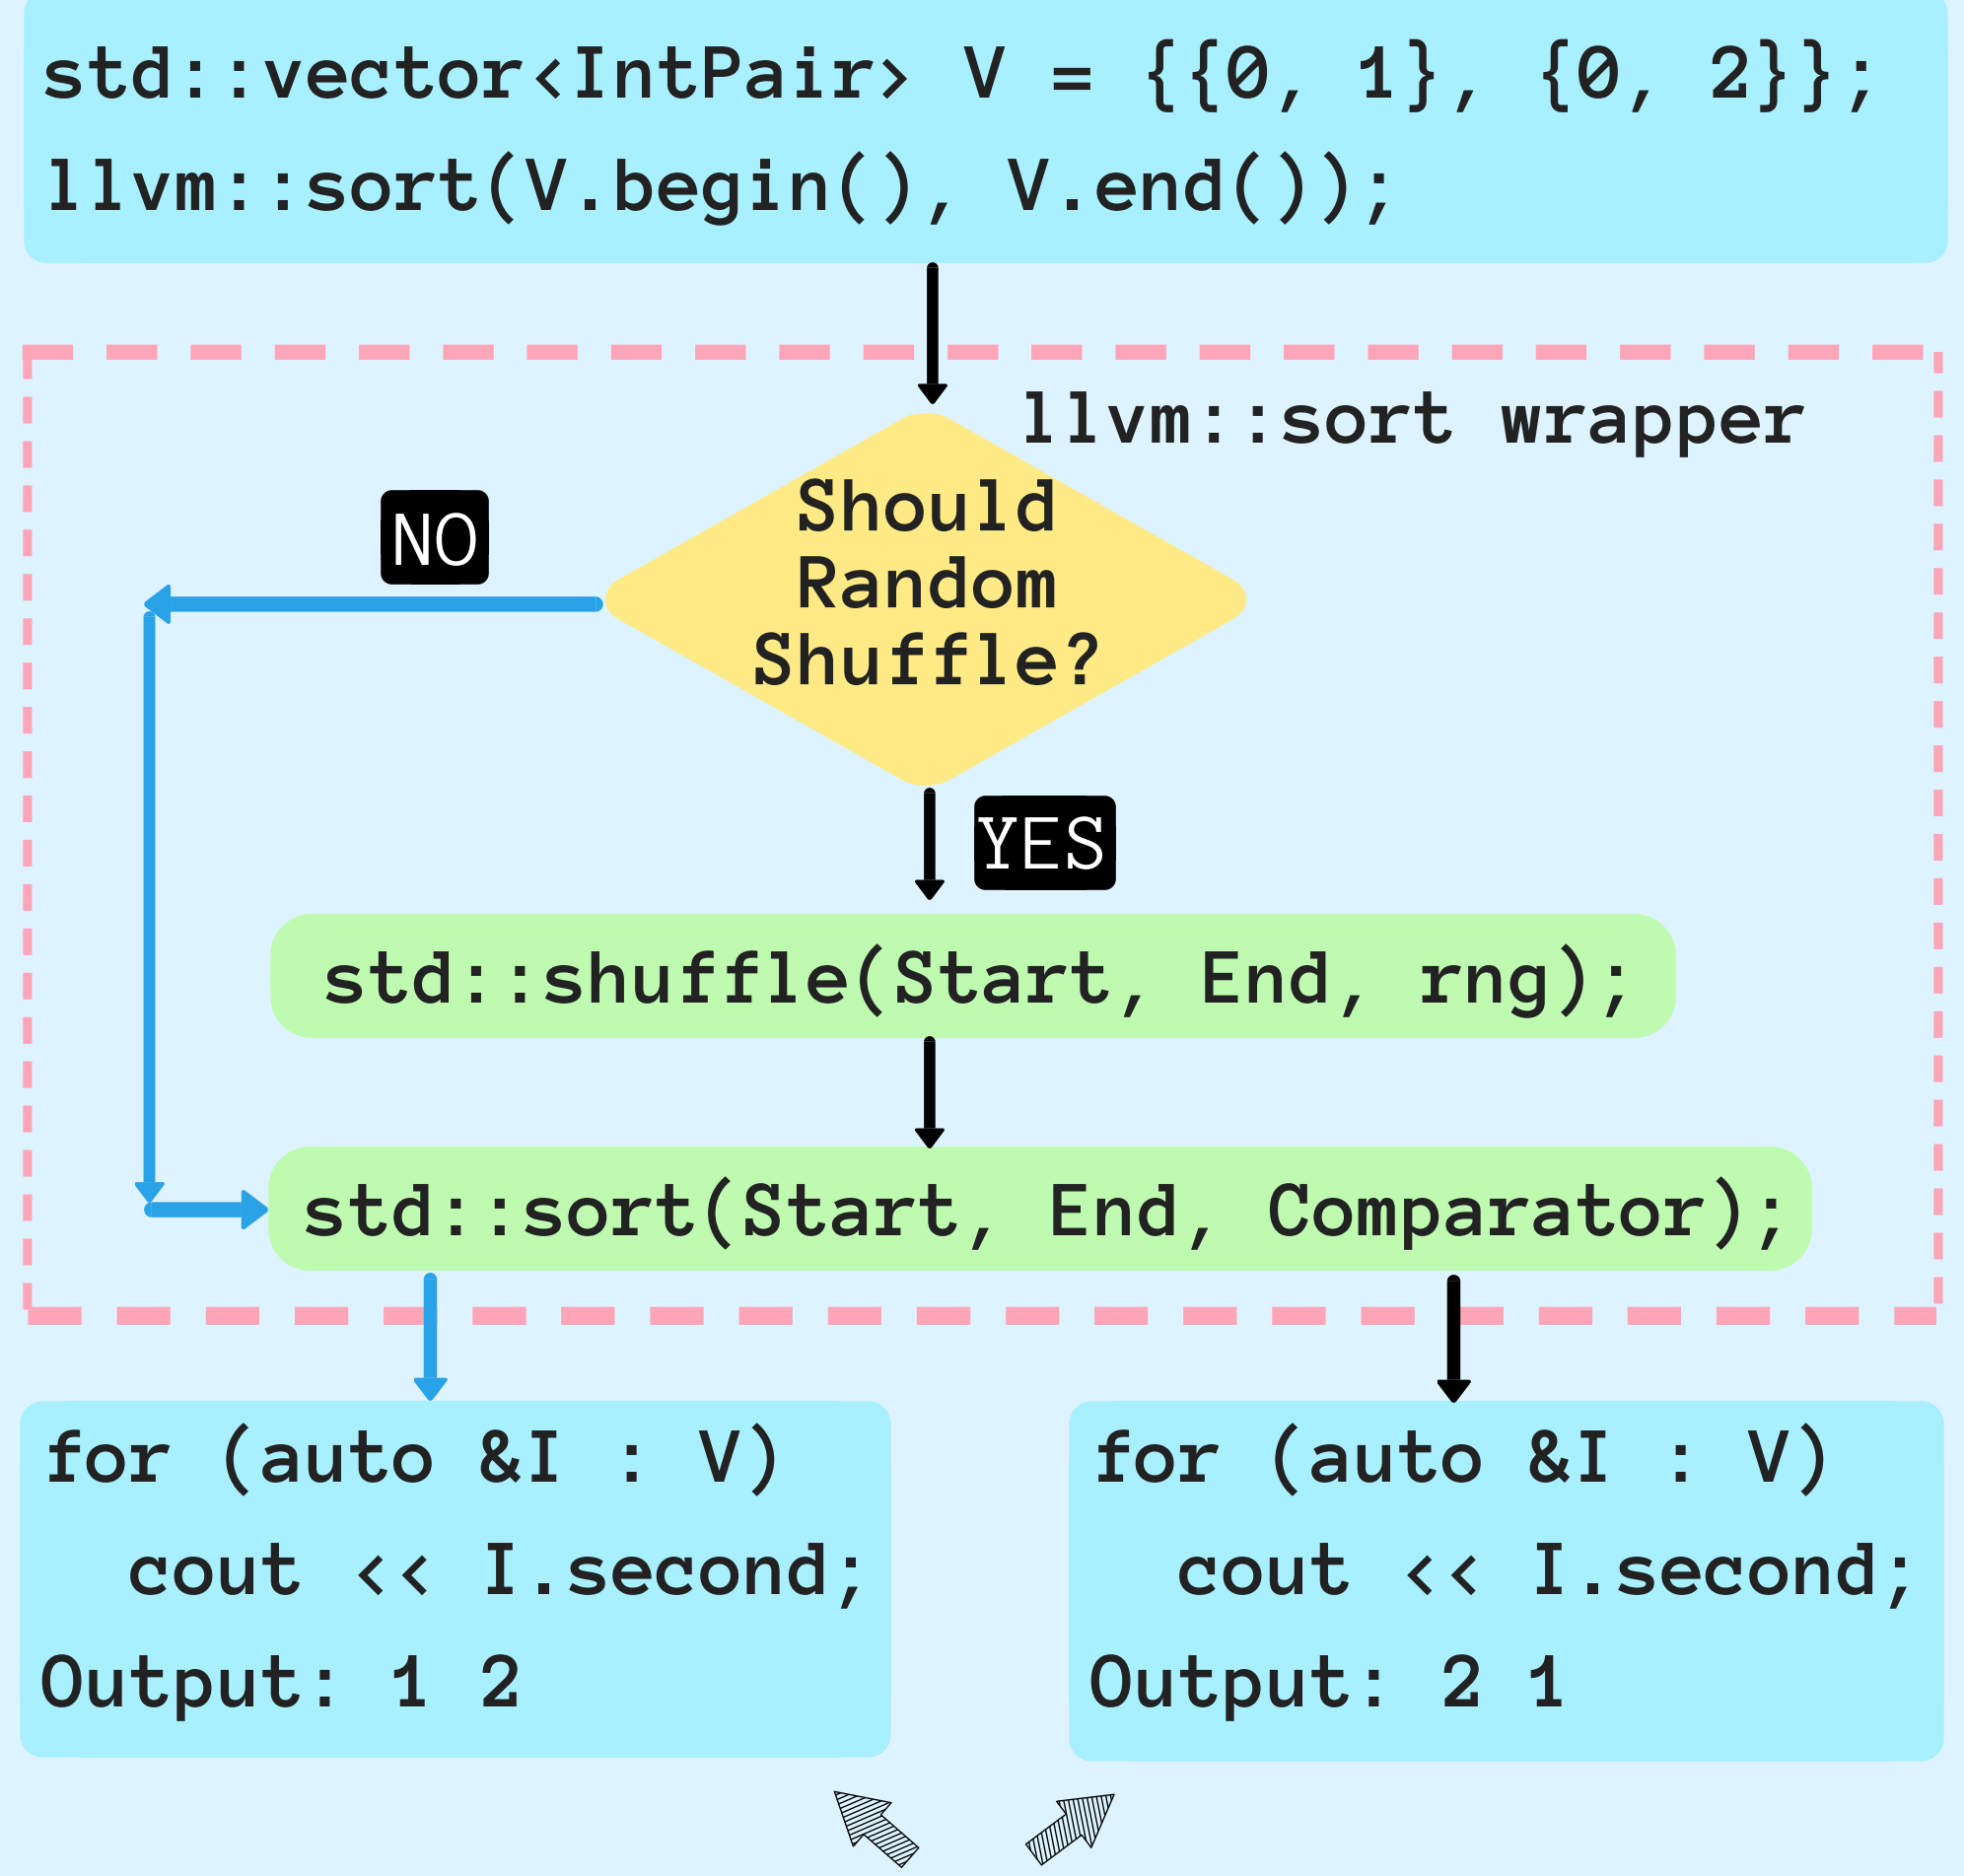


Output sensitive to iteration order.
Possible non-determinism!

6. HOW DOES LLVM DETECT NON-DETERMINISTIC SORTING ORDER?

- Added a wrapper `llvm::sort` which randomly shuffles a container before invoking `std::sort`.
- Uncovers non-deterministic sorting of keys with same values.
- Bugs uncovered and fixed:

44



Different sorting order.
Possible non-determinism!

7. HOW DO YOU FIX/AVOID NON-DETERMINISM?

i

Use an ordered container

```
std::vector<T> V;
for (auto I : V)
```

iii

Use a stronger sort predicate

```
std::sort(V.begin(), V.end(),
[] (T a, T b) { return
a.first < b.first &&
a.second < b.second; });
```

ii

Sort the container before iteration

```
std::sort(V.begin(), V.end());
for (auto I : V)
```

iv

Use a stable sort function

```
std::stable_sort(V.begin(), V.end());
```

8. WHERE CAN I LEARN MORE?

<https://github.com/mgrang/non-determinism>

