


NON-DETERMINISM IN LLVM CODE GENERATION

Mandeep Singh Grang (Engineer, Senior) • Qualcomm Innovation Center, Inc., San Diego, CA

1. THE PROBLEM

LLVM generates non-deterministic code



2. WHERE IS THE PROBLEM OBSERVED?

- Between back-to-back runs of the same toolchain
- Between asserts and non-asserts toolchains
- Between toolchains hosted on different operating systems

3. WHAT CAUSES THIS NON-DETERMINISM?

i

Iteration of unordered containers

```
int x = 1, y = 2, z = 3, w = 4;
```

```
DenseMap<int *, int> Map;  
Map[&x] = x;  
Map[&y] = y;  
Map[&z] = z;  
Map[&w] = w;
```

```
for (auto &I : Map)  
  dbgs() << I.second;
```

Output of 3 runs:

```
1 2 3 4  
2 1 4 3  
2 3 4 1
```

The DenseMap hashes on pointer keys which are different from run-to-run. This makes the iteration order non-deterministic.

ii

Using non-stable sorting functions

```
using IntPair = std::pair<int, int>;  
IntPair x = {1, 1}, y = {1, 2},  
        z = {1, 3}, w = {1, 4};  
SmallVector<IntPair, 4> V = { x, y, z, w };
```

```
std::sort(V.begin(), V.end(),  
[] (IntPair a, IntPair b) {  
    return a.first < b.first; });
```

```
for (auto I : V)  
  dbgs() << I.second;
```

Sorting order of keys with the same values is non-deterministic.

Output: Depends on the underlying std::sort algorithm, which could be non-stable.

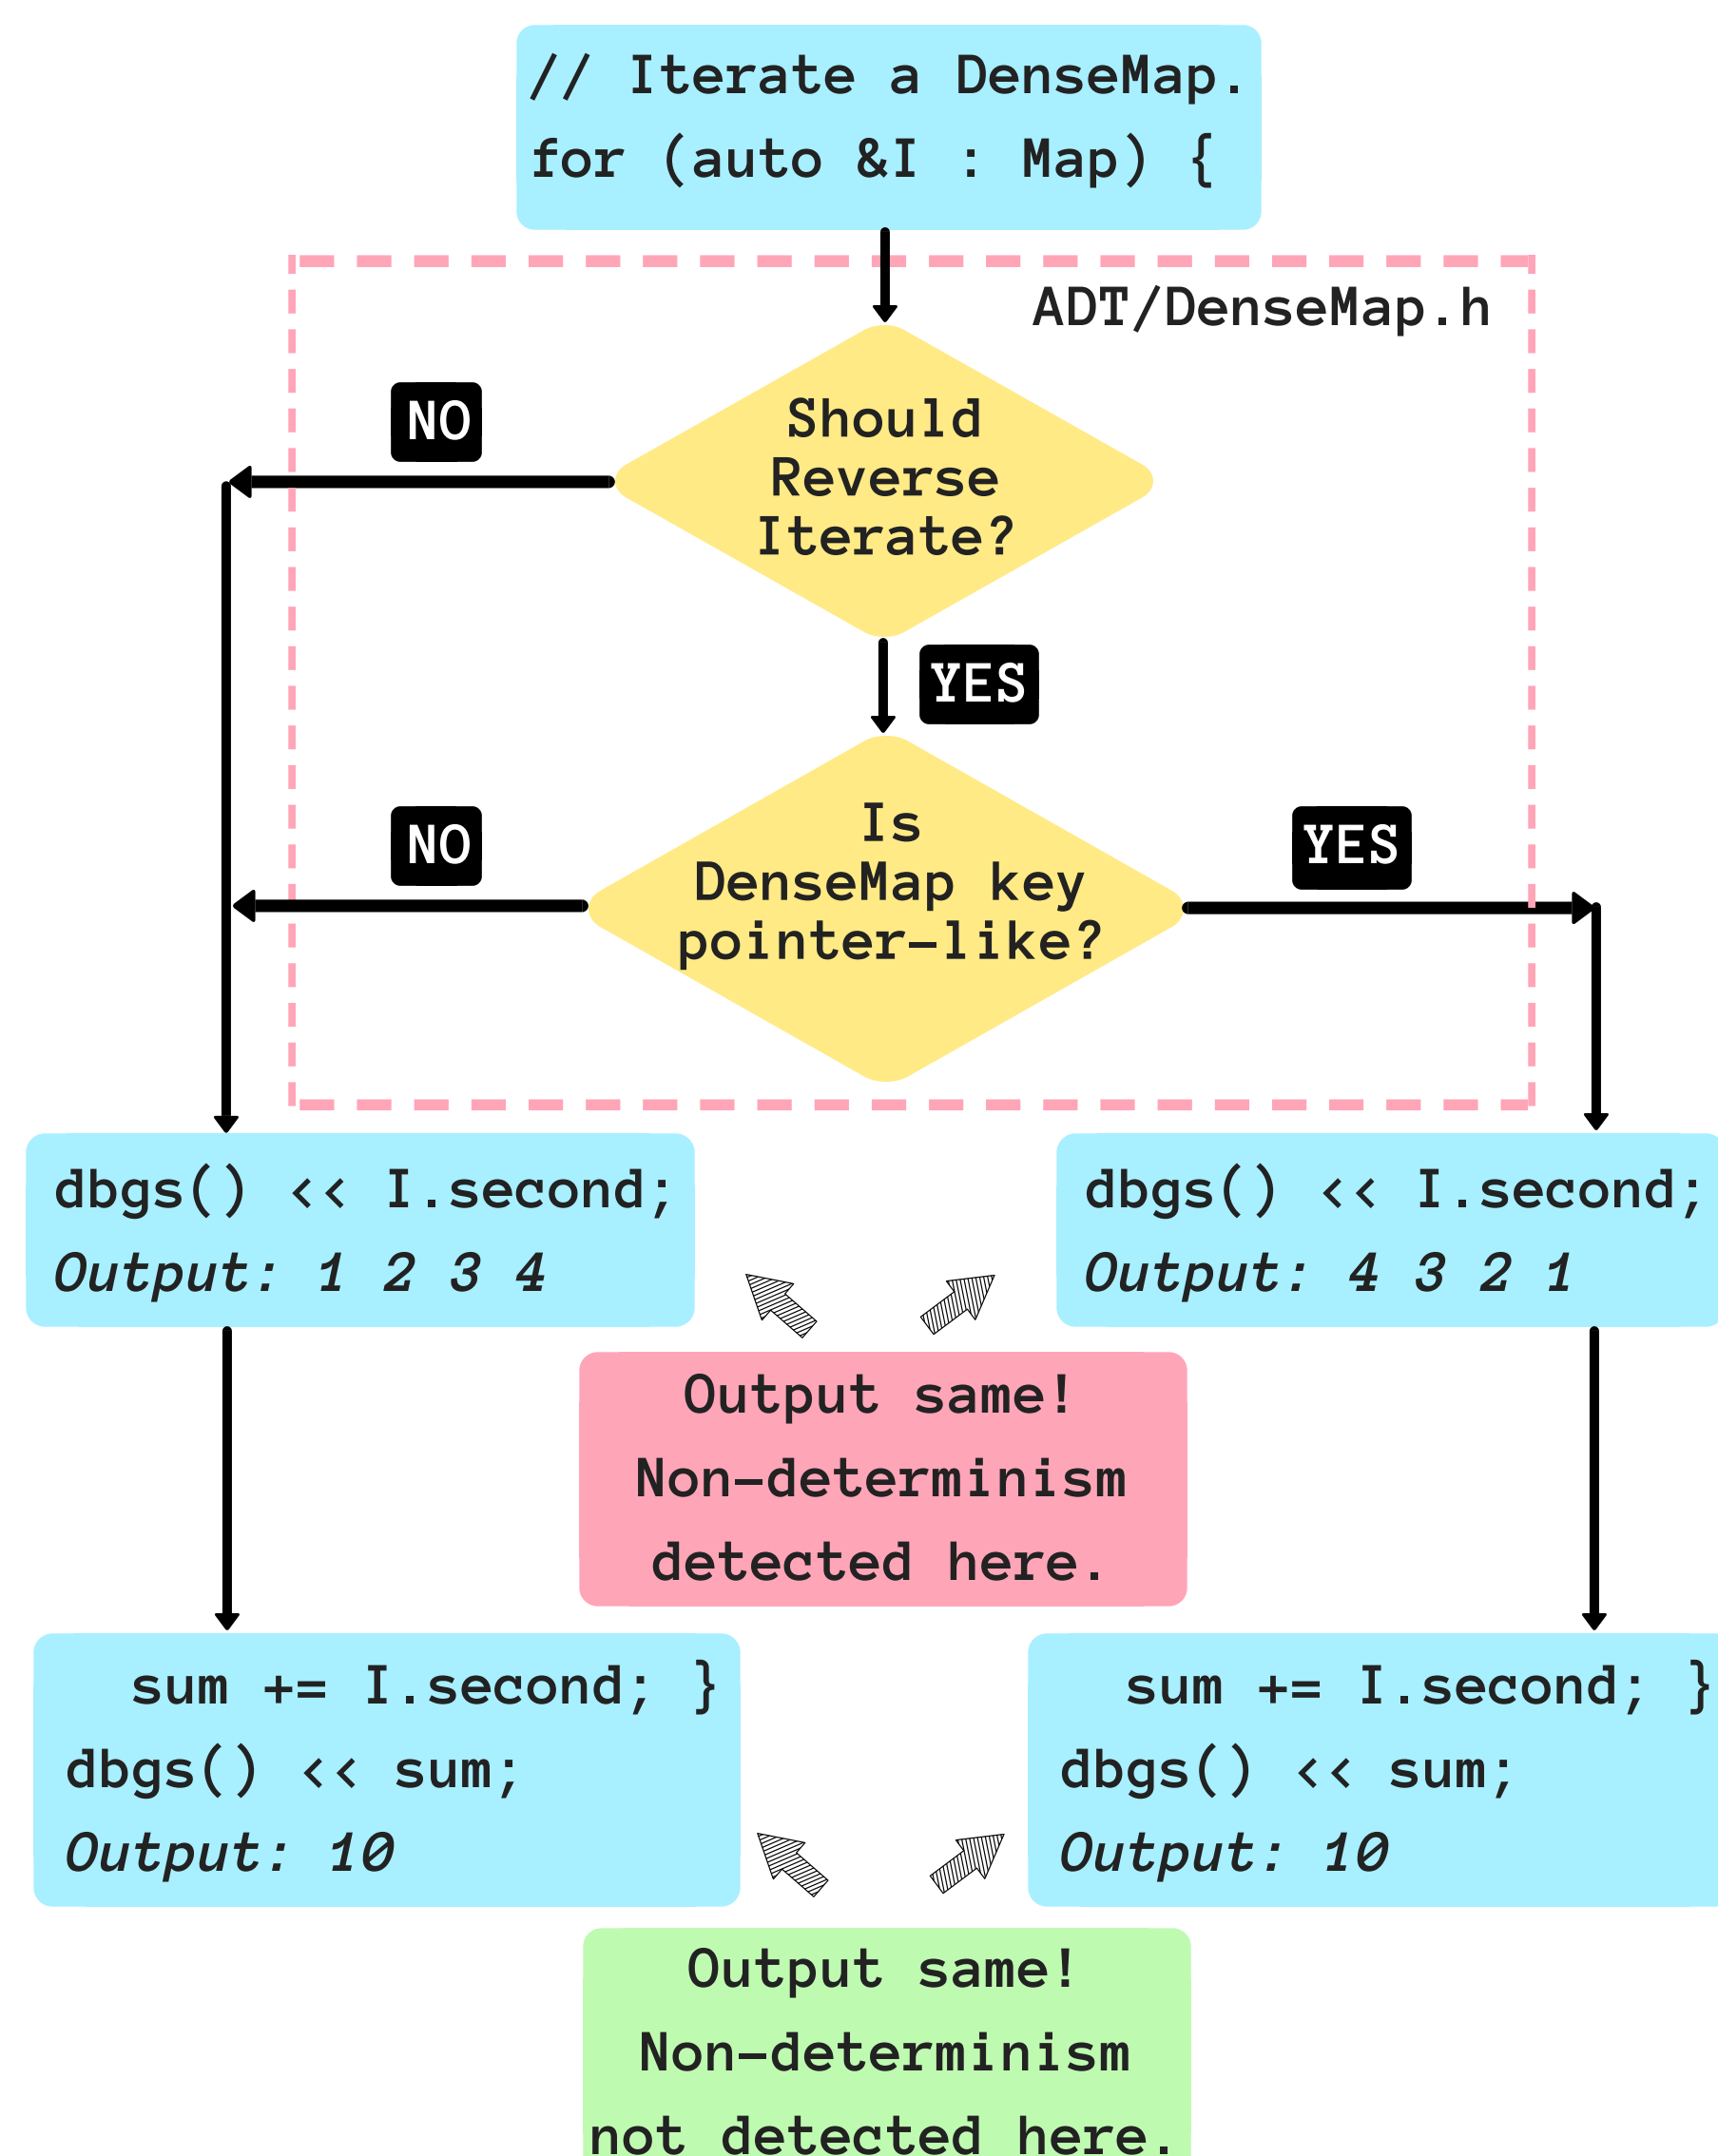
4. WHY IS IT A PROBLEM?

The generated code might not necessarily be "wrong". However, it may cause:

- Hard-to-reproduce bugs
- Unexpected runtime crashes
- Unpredictable performance

5. HOW DO YOU UNCOVER INSTANCES OF NON-DETERMINISM?

- The idea is to force "reverse iteration" of unordered containers and check if the generated code changes.
- Reverse iteration is limited to containers with pointer-like keys.



- Currently, the following containers support reverse iteration:

SmallPtrSet, DenseMap and DenseSet

- To enable reverse iteration by default, set the build macro:

-DLLVM_REVERSE_ITERATION:BOOL=ON

- A "Reverse Iteration" buildbot has been setup:

<http://lab.llvm.org:8011/builders/reverse-iteration>

6. HOW DO YOU FIX/AVOID NON-DETERMINISM?

i

Sort the container before iteration

```
SmallPtrSet<T *, 4> S;  
std::sort(S.begin(), S.end(), [] (T *A, T *B)  
{ return A->x < B->x; });  
for (auto &I : S);
```

ii

Use a stronger sort predicate

```
std::sort(S.begin(), S.end(), [] (T *A, T *B)  
{ return A->x < B->x && A->y < B->y; });
```

iii

Use a stable sort function

```
std::stable_sort(S.begin(), S.end(),  
[] (T *A, T *B) { return A->x < B->x; });
```

iv

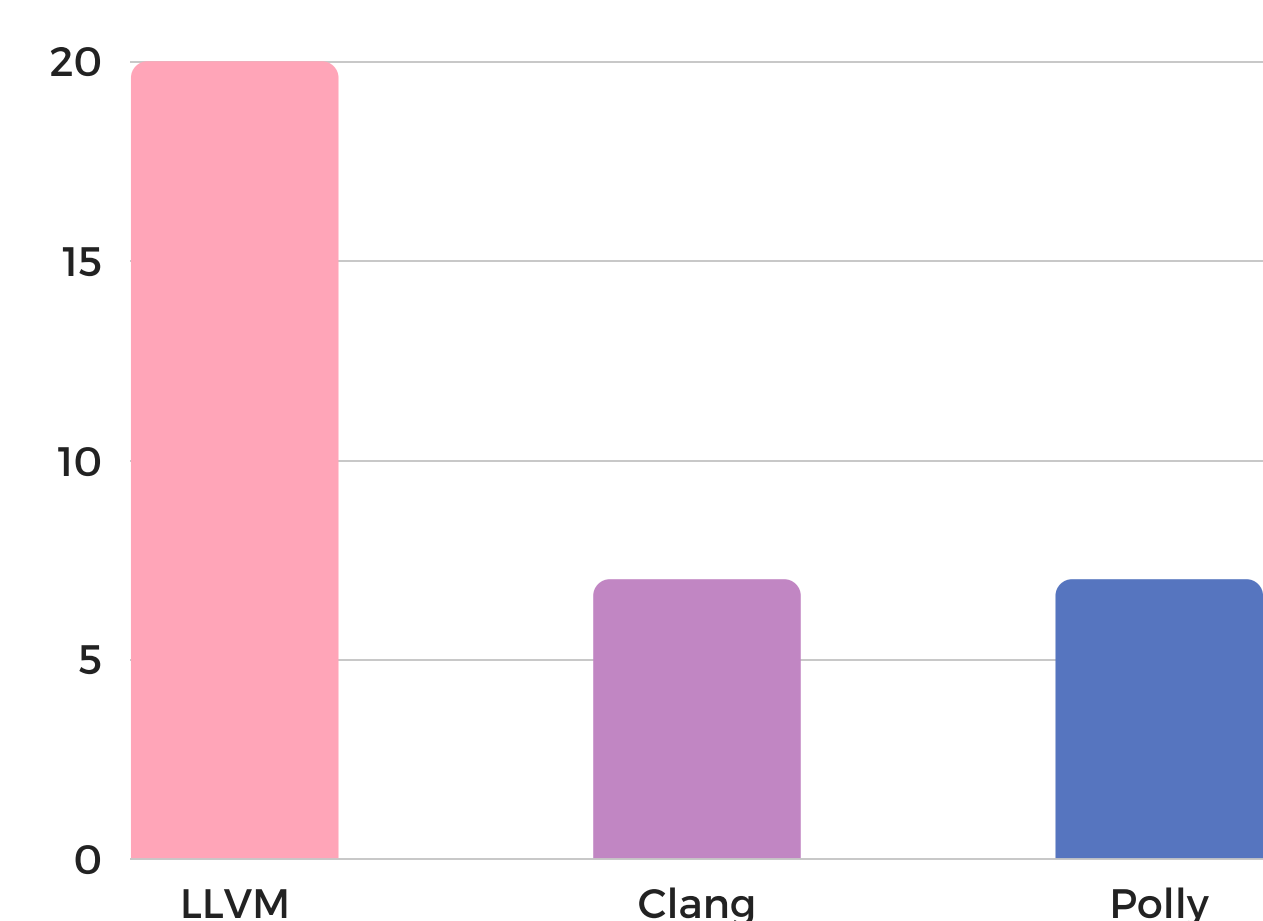
Use an ordered container

```
SmallSetVector<T *, 4> S;  
for (auto &I : S);
```



Using ordered containers instead of unordered ones can increase compile time.

7. HOW MANY ISSUES HAVE BEEN DETECTED SO FAR?



ninja check failures detected by reverse iteration
(Oct 2016 - Oct 2017)